

# CS 240 Homework 2

Wei Yuxiang ID: 206529292

## Q1. Prove that Parity $\in \text{NC}^1$

We show that  $\text{Parity}(x_1, \dots, x_n) = x_1 \oplus x_2 \oplus \dots \oplus x_n$  can be computed by a Boolean circuit of depth  $O(\log n)$  and size  $O(n)$ .

Each XOR gate can be expressed as

$$a \oplus b = (a \vee b) \wedge \neg(a \wedge b),$$

which is a constant-depth circuit using {AND, OR, NOT} gates.

Arrange the XOR gates in a balanced binary tree: each internal node computes the XOR of its two children, and the root outputs the overall parity. Hence,

$$\text{depth} = \lceil \log_2 n \rceil = O(\log n), \quad \text{size} = n - 1 = O(n).$$

Since a log-space machine can describe this circuit family, it is uniform. Therefore,

$$\text{PARITY} \in \text{NC}^1.$$

## Q2. Prove that $\text{NC}^i \subseteq \text{AC}^i \subseteq \text{NC}^{i+1}$

**Definitions.** A language is in  $\text{NC}^i$  if it can be computed by a family of Boolean circuits of polynomial size, depth  $O(\log^i n)$ , and bounded fan-in (at most 2). The class  $\text{AC}^i$  is defined similarly but allows *unbounded* fan-in for AND and OR gates.

(1)  $\text{NC}^i \subseteq \text{AC}^i$ . Every bounded fan-in gate is a special case of an unbounded one. Hence, any  $\text{NC}^i$  circuit is also an  $\text{AC}^i$  circuit with the same depth and size. Therefore  $\text{NC}^i \subseteq \text{AC}^i$ .

(2)  $\text{AC}^i \subseteq \text{NC}^{i+1}$ . Conversely, consider an  $\text{AC}^i$  circuit of depth  $O(\log^i n)$ . Each unbounded fan-in gate (say an AND or OR with  $k$  inputs) can be replaced by a balanced binary tree of the same gate type, where each node has fan-in 2. This increases the depth by  $O(\log k)$ . Since  $k \leq n^{O(1)}$ , we have  $O(\log k) = O(\log n)$ . Thus the overall circuit depth becomes

$$O(\log^i n) + O(\log n) = O(\log^{i+1} n),$$

and the size remains polynomial.

**Conclusion.** Replacing each unbounded gate by a binary tree increases the depth by at most one logarithmic factor. Hence any  $\text{AC}^i$  circuit can be simulated by an  $\text{NC}^{i+1}$  circuit. Together we have

$$\text{NC}^i \subseteq \text{AC}^i \subseteq \text{NC}^{i+1}.$$

### Q3. Prove that the evaluation of FO formulae is in $\text{AC}^0$

We consider the *data complexity* of evaluating a fixed first-order (FO) sentence on a finite structure. Let  $\text{dom}$  be the domain of size  $n$ . An FO sentence is built from atomic predicates, Boolean connectives, and quantifiers  $\exists, \forall$ .

**Translation to circuits.** Each atomic predicate  $R(x_1, \dots, x_k)$  is evaluated by reading the corresponding entry in the input relation  $R$ . Boolean connectives correspond to standard logic gates:

$$\neg\varphi \mapsto \text{NOT}(\varphi), \quad \varphi_1 \wedge \varphi_2 \mapsto \text{AND}(\varphi_1, \varphi_2), \quad \varphi_1 \vee \varphi_2 \mapsto \text{OR}(\varphi_1, \varphi_2).$$

Quantifiers are interpreted as unbounded fan-in gates over all elements of the domain:

$$\exists x \varphi(x) \mapsto \bigvee_{a \in \text{dom}} \varphi(a), \quad \forall x \varphi(x) \mapsto \bigwedge_{a \in \text{dom}} \varphi(a).$$

**Depth and size.** If the FO sentence has  $k$  nested quantifiers, the resulting circuit has depth

$$\text{depth} = O(k),$$

since each quantifier adds one layer of unbounded gates. Because the formula is fixed,  $k$  is a constant, so  $\text{depth} = O(1)$ . Each quantifier ranges over  $n$  elements, giving polynomially many gates:

$$\text{size} = n^{O(1)}.$$

**Conclusion.** The circuit family evaluating the FO sentence has constant depth, polynomial size, and unbounded fan-in AND/OR gates. Therefore,

$$\text{FO evaluation} \in \text{AC}^0.$$

### Q4. JOB 33c as a Conjunctive Query (join-only)

We ignore all non-join predicates and drop attributes not used in joins. Using shared variables to encode equality joins, a boolean CQ for JOB 33c is:

$$\begin{aligned} Q() :- & \text{company\_name}(c_1), \text{company\_name}(c_2), \\ & \text{info\_type}(it_1), \text{info\_type}(it_2), \\ & \text{kind\_type}(k_1), \text{kind\_type}(k_2), \\ & \text{link\_type}(\ell), \\ & \text{title}(m_1, k_1), \text{title}(m_2, k_2), \\ & \text{movie\_link}(\ell, m_1, m_2), \\ & \text{movie\_info\_idx}(it_1, m_1), \text{movie\_info\_idx}(it_2, m_2), \\ & \text{movie\_companies}(c_1, m_1), \text{movie\_companies}(c_2, m_2). \end{aligned}$$

*Variable meaning:*  $m_1, t_1$ -side movie id;  $m_2, t_2$ -side linked movie id;  $k_1, k_2$  kind ids;  $it_1, it_2$  info\_type ids;  $c_1, c_2$  company ids;  $\ell$  link\_type id.

## Q5: Query Hypergraph for 33c

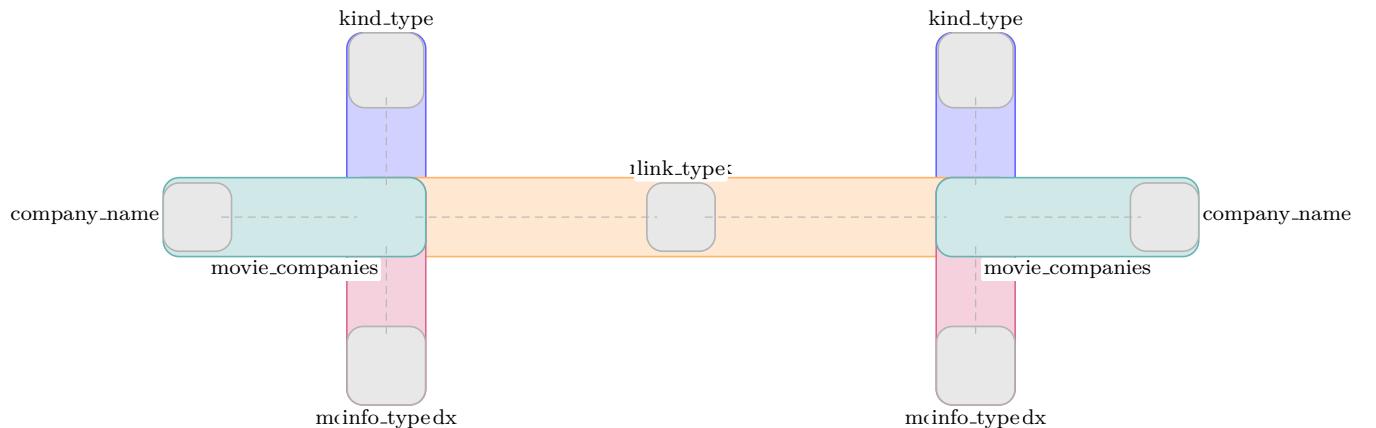
We use the conjunctive query (from Problem 4) that keeps only join predicates and join-participating attributes. Let the variable set (hypergraph vertices) be

$$V = \{ c_1, c_2, it_1, it_2, k_1, k_2, \ell, m_1, m_2 \}.$$

Each relation atom induces a hyperedge on the variables it mentions. Thus the hyperedge set is

$$\begin{aligned} E = \{ & \{c_1\}_{\text{company\_name}}, \{c_2\}_{\text{company\_name}}, \{it_1\}_{\text{info\_type}}, \{it_2\}_{\text{info\_type}}, \\ & \{k_1\}_{\text{kind\_type}}, \{k_2\}_{\text{kind\_type}}, \{\ell\}_{\text{link\_type}}, \\ & \{m_1, k_1\}_{\text{title}}, \{m_2, k_2\}_{\text{title}}, \\ & \{\ell, m_1, m_2\}_{\text{movie\_link}}, \\ & \{it_1, m_1\}_{\text{movie\_info\_idx}}, \{it_2, m_2\}_{\text{movie\_info\_idx}}, \\ & \{c_1, m_1\}_{\text{movie\_companies}}, \{c_2, m_2\}_{\text{movie\_companies}} \}. \end{aligned}$$

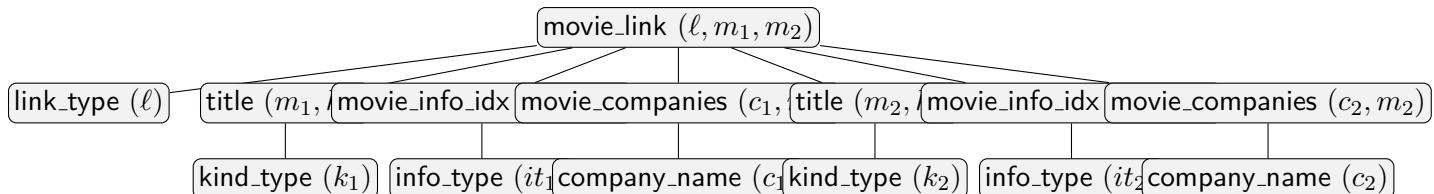
**TikZ drawing of the hypergraph.**



*Explanation.* Each vertex represents a join variable shared across relations. Each colored region denotes one relation (a hyperedge) covering the variables that appear in it. The central edge  $\text{movie\_link}(\ell, m_1, m_2)$  connects the two symmetric parts of the query through movies  $m_1$  and  $m_2$ .

## Q6: A Join Tree for 33c (rooted at **movie\_link**)

A valid join tree that satisfies the running–intersection property can be rooted at  $\text{movie\_link}(\ell, m_1, m_2)$  and attach the atoms that share variables with it. Unaries then hang from their respective binary atoms.



*This join tree satisfies the running-intersection property, with  $\text{movie\_link}(\ell, m_1, m_2)$  as the root.*

## Q7: Implement Query 33c with Yannakakis' Algorithm

Implemented Yannakakis' algorithm for query 33c using the join tree rooted at `movie_link`. The left branch includes  $\{t_1, it_1, mi\_idx1, mc_1, cn_1, kt_1\}$ , and the right branch includes  $\{t_2, it_2, mi\_idx2, mc_2\}$ . All join and non-join predicates are included.

The algorithm follows three steps: (1) bottom-up semijoin reduction, (2) top-down semi-join reduction, (3) final join and aggregation. Non-join filters (e.g., `rating < 3.5`) are handled by DuckDB during data loading.

The result was compared with the SQL baseline run directly in DuckDB. Both outputs match exactly, each returning one row with all NULL values, confirming correctness.

```
[SQL baseline output]
first_company second_company first_rating second_rating first_movie second_movie
0           None          None        None       None      None      None

[Yannakakis output]
first_company second_company first_rating second_rating first_movie second_movie
0           None          None        None       None      None      None

[Compare] Yannakakis matches SQL baseline? True
```

Figure 1: Comparison between SQL baseline and Yannakakis algorithm outputs. Both return a single row with all NULL values.

**Reproducibility.** I provide a single script `src/hw2_q7_yannakakis_33c.py`. It reads the original JOB CSVs from `./data` by default (configurable via `--data-dir`). No external DB is required; I use DuckDB's in-memory CSV reader. Running `python3 src/hw2_q7_yannakakis_33c.py` prints the SQL baseline result, our Yannakakis result, and a boolean equality check.